

DATA ANALYTICS PORTFOLIO PROJECT

Global Superstore Sales & Profit Analysis

A Comprehensive Diagnostic Analysis of Revenue, Profitability, Customer Behaviour and Regional Performance (2011–2014)

Author	Behnia Shayeste
Role	Data Analyst / Data Scientist
Project Type	Data Analytics Portfolio Project
Tools Used	Python, Pandas, NumPy, Matplotlib, Jupyter Notebook
Analysis Period	2011 – 2014
Document Type	Final Business Report

This report extends the original analytical findings with additional senior-level interpretation, statistical framing, risk assessment, and an implementation roadmap intended for executive and analytics-team audiences.

Table of Contents

Table of Contents	2
1. Executive Summary	3
2. Project Objectives	3
3. Dataset Overview	3
Engineered Features	4
4. Methodology	4
5. Key Findings	4
5.1 Sales Performance	4
5.2 Profit Performance	5
5.3 Customer Analysis	5
5.4 Product Performance	5
5.5 Regional Performance	5
5.6 Discount Analysis	5
6. Statistical Notes & Analytical Rigor	6
7. Risk Assessment	6
8. Business Insights	6
9. Business Recommendations	7
10. Suggested A/B Tests	7
11. Project Limitations	8
12. Future Work	8
13. Conclusion	9

1. Executive Summary

This report presents an end-to-end diagnostic analysis of the Global Superstore dataset, covering approximately 25,000 orders and 1,590 customers across the 2011–2014 period. The objective was to move beyond descriptive reporting and identify the structural drivers of revenue, profitability, and customer value, using a reproducible Python/Pandas workflow encompassing data cleaning, feature engineering, exploratory data analysis (EDA), and business-oriented synthesis.

The headline finding is that Technology is the company's strongest category on both revenue and profit dimensions, while Furniture exhibits a structural profitability problem: strong topline sales are not converting into proportional margin. Discounting emerges as the single largest controllable lever on profitability — roughly a quarter of all orders are loss-making, and this loss rate is heavily concentrated in the Medium and High discount bands. Customer concentration risk is low (revenue is broadly distributed across the customer base), but geographic concentration risk is material, with a small set of countries consistently generating negative margins.

Senior-analyst framing: profitability in this dataset is driven less by what is sold and more by how it is sold — discount depth and geography explain a disproportionate share of order-level profit variance relative to product category alone. This reframes the recommended interventions toward tighter discount governance layered on top of category-level investment.

Total Sales	Total Profit	Orders	Customers	Loss-Order Rate
\$12.64M	\$1.47M	25,035	1,590	24.5%

2. Project Objectives

- Quantify overall sales and profitability trends across the 2011–2014 period.
- Identify top-performing and underperforming products and sub-categories.
- Characterize customer purchasing behaviour and assess revenue concentration risk.
- Evaluate regional, market, and country-level performance, including loss-making geographies.
- Isolate the impact of discount strategy on order-level and category-level profitability.
- Translate analytical findings into prioritized, testable business recommendations.

3. Dataset Overview

The dataset spans four fiscal years and covers global order-level transactions, joined with customer, product, and shipping attributes.

Metric	Value
Total Sales	\$12.64 Million
Total Profit	\$1.47 Million
Implied Overall Margin	≈ 11.6%
Orders	25,035
Customers	1,590

Metric	Value
Avg. Revenue per Customer	≈ \$7,951
Analysis Period	2011 – 2014

Engineered Features

- **Shipping Days** — elapsed time between order and ship date, used as a logistics-quality proxy.
- **Profit Margin** — profit / sales at the order line level, the primary profitability normalisation metric.
- **Sales per Unit** — unit economics indicator, useful for identifying mispriced SKUs.
- **Discount Category** — banded discount levels (None / Low / Medium / High) used throughout the discount analysis.
- **Quarter** — calendar quarter, enabling seasonality decomposition.
- **Loss Order Flag** — binary indicator (profit < 0), the basis for the 24.5% loss-rate statistic.

4. Methodology

The project followed a standard, reproducible data analytics workflow, structured to separate data preparation from inference and to keep each stage independently auditable:

- **Data Understanding** — schema review, type validation, missingness and duplicate checks.
- **Data Cleaning** — handling of nulls, type coercion, outlier review on Sales/Profit/Discount.
- **Feature Engineering** — construction of the derived fields listed in Section 3.
- **Exploratory Data Analysis (EDA)** — univariate and bivariate analysis across time, product, customer, and geography.
- **Business Analysis** — translation of statistical patterns into operational implications.
- **Business Recommendations** — prioritized, testable interventions (Sections 9–10).

The exploratory phase was organized around seven analytical lenses: KPI Analysis, Sales Analysis, Profit Analysis, Customer Analysis, Product Analysis, Regional Analysis, and Shipping Analysis — each cross-referenced against discount level and time period to avoid drawing conclusions from a single dimension in isolation.

5. Key Findings

5.1 Sales Performance

Annual sales grew consistently across the analysis window, peaking in 2014. Quarterly decomposition shows a pronounced Q4 effect, with November and December as the two strongest months — a pattern consistent with holiday-season retail demand and one that should anchor annual inventory and marketing planning cycles.

Senior-analyst note: because the Q4 effect is consistent across all four years, it should be treated as a structural seasonal pattern rather than a one-off anomaly, and built directly into demand-forecasting baselines (e.g. seasonal-naïve or SARIMA models) rather than smoothed out.

5.2 Profit Performance

Profit tracks the same upward trajectory as sales, but the two are not proportionally linked at the category level. Technology generates the highest total profit of any category; Furniture generates high revenue but disproportionately low profit. Approximately 24.5% of all orders are loss-making, and average profit declines sharply and monotonically as discount depth increases — Medium and High discount bands show negative average profit.

5.3 Customer Analysis

The Consumer segment leads on every primary metric — sales, profit, customer count, and order count — making it the structural backbone of the business rather than a niche segment. Average revenue per customer is approximately \$7,951.

A Pareto (80/20) analysis shows that approximately 42.5% of customers account for 80% of total sales. This is materially more diffuse than a classic 80/20 distribution, indicating low customer concentration risk: the business is not dependent on a small set of key accounts, which is favourable from a revenue-stability and churn-resilience standpoint.

5.4 Product Performance

Technology is the strongest-performing category overall. Within it, Phones and Copiers stand out, alongside Chairs and Bookcases from the Furniture category as comparatively strong sub-category performers. Several smartphone SKUs are consistently among the highest-selling products, and the Canon imageCLASS 2200 Advanced Copier is the single highest-profit product in the catalogue.

Conversely, a recurring set of Tables and Printers generate substantial financial losses, and Furniture as a whole shows low profit margins despite healthy sales volume — reinforcing the category-level finding in 5.2 and suggesting the issue is structural (pricing, logistics cost, or discount exposure) rather than driven by a handful of outlier SKUs.

5.5 Regional Performance

The Central region leads on total sales. At the market level, APAC generates the highest revenue of any global market, while the United States remains the largest single-country market and leads on both sales and profit simultaneously — making it the most balanced high-performing market in the dataset.

A consistent group of countries — Turkey, Nigeria, Netherlands, Honduras, and Pakistan — generate persistent financial losses, and margin analysis shows negative average profit margins for the Africa and EMEA markets overall. The persistence of these losses across the analysis period (rather than a single bad year) suggests structural causes — likely shipping cost, local discount practice, or product mix — rather than transient market shocks.

5.6 Discount Analysis

Discount strategy is the most powerful lever identified in this analysis. Orders with no discount generate the highest average profit; Medium and High discount bands generate negative average profit; and higher discount depth is strongly associated with loss-making transactions across virtually every category and region examined. This pattern is consistent enough to be treated as a near-universal rule in the current dataset rather than a category-specific exception, and is the clearest candidate for a formal policy change (see Recommendation 1 and Test 1).

6. Statistical Notes & Analytical Rigor

This section adds methodological context that a senior data scientist would typically flag before findings are operationalized, to ensure the recommendations in Sections 9–10 are read with appropriate statistical caution.

- **Correlation vs. causation:** the discount–profit relationship is strongly correlational in this analysis. High discounts are also more likely to be applied to already low-margin or clearance-style products, so part of the observed effect may reflect product selection rather than discount depth alone. A regression or matched-sample analysis controlling for product and category would isolate the discount effect more cleanly before the policy is finalized.
- **Country-level loss attribution:** the loss-making countries identified in 5.5 should be decomposed by discount level, shipping cost, and currency/price-list effects before any market-exit or pricing decision is made — a country can appear unprofitable due to a small number of high-discount enterprise orders rather than systemic market weakness.
- **Pareto threshold sensitivity:** the 42.5% figure in 5.3 should be re-validated on a rolling basis (e.g. quarterly) since customer concentration can drift as the customer base grows; a single point-in-time Pareto curve can understate emerging concentration risk.
- **Seasonality vs. trend conflation:** the year-over-year sales growth noted in 5.1 should be decomposed into trend and seasonal components (e.g. via STL decomposition) before being used as an input to growth targets, to avoid mistaking a strong Q4 for sustained underlying growth.

7. Risk Assessment

Risk	Likelihood	Impact	Mitigation
Discount policy change reduces order volume	Medium	Medium	Pilot via A/B test (Test 1) on a discount-threshold segment before full rollout.
Furniture rationalisation reduces basket completeness	Low	Medium	Review bundled / cross-sell impact before discontinuing SKUs.
Exiting loss-making countries forfeits growth options	Medium	Medium	Diagnose root cause (5.5/6) before deciding between repricing and exit.
Incorrect Q4 demand forecast causes stock-outs/overstock	Medium	High	Use seasonally-decomposed forecasts with safety-stock buffers ahead of Q4.

8. Business Insights

- Technology is the clearest growth lever, combining high revenue with high, durable profitability.
- Furniture has a structural margin problem: strong sales are not converting into proportional profit and warrant a pricing/portfolio review rather than a sales-volume push.
- Discount depth is the single largest controllable driver of order-level losses across the dataset.
- Q4 (especially Nov–Dec) is a reliable, multi-year seasonal peak that should anchor inventory and marketing calendars.

- Customer concentration risk is low — the business is not dependent on a small group of accounts.
- A small, consistent set of countries (Turkey, Nigeria, Netherlands, Honduras, Pakistan) require individual, root-cause-level investigation rather than a blanket regional policy.

9. Business Recommendations

Recommendations are presented in priority order, weighted by estimated profitability impact relative to implementation effort.

#	Recommendation	Rationale	Priority
1	Review and tighten the discount policy, capping or restructuring high-discount approvals.	Medium/High discount bands show negative average profit and drive a large share of loss-making orders.	High
2	Increase investment in Technology (marketing spend, assortment, inventory depth).	Highest combination of revenue and profit margin of any category.	High
3	Re-evaluate Furniture pricing and SKU portfolio, focusing on Tables and Printers.	Persistent low/negative profitability despite strong sales volume.	Medium
4	Conduct root-cause investigation into loss-making countries (Turkey, Nigeria, Netherlands, Honduras, Pakistan).	Losses are persistent and structural, not one-off; cause (cost, discount, mix) is not yet isolated.	Medium
5	Build Q4-specific inventory and marketing plans ahead of the November–December peak.	Multi-year, statistically consistent seasonal pattern.	Medium
6	Develop a retention/loyalty program for the Consumer segment.	Largest contributor to both sales and profit; low concentration risk makes broad retention efficient.	Low–Med

10. Suggested A/B Tests

Each recommendation above should be validated experimentally before full rollout. Proposed test designs, primary KPIs, and suggested minimum run durations are summarized below.

Test	Hypothesis	Primary KPI(s)	Duration
1	Reducing discounts above 30% increases average profit per order.	Average Profit per Order	6–8 weeks
2	Increased promotional spend on Technology products increases total revenue.	Sales; Profit	1 quarter
3	Removing low-performing Furniture SKUs improves category-level profitability.	Profit Margin (category)	1 quarter
4	Pre-Q4 marketing campaigns increase seasonal revenue capture.	Quarterly Sales; Conversion Rate	8–10 weeks

Implementation note: tests should be randomized at the customer or order level where feasible (not the country/region level) to avoid confounding from macro-economic or currency effects already observed in the regional analysis (Section 5.5).

11. Project Limitations

- The analysis is based solely on historical transactional data; no external variables (competitor pricing, inflation, macroeconomic conditions, marketing spend) were available to control for confounding effects.
- Discount-profit relationships are correlational, not causal — see Section 6 for the specific caveats that should be resolved before the discount policy (Recommendation 1) is finalized.
- Suggested A/B tests (Section 10) are design proposals only and have not yet been implemented or run.
- Country-level loss findings are aggregate and have not yet been decomposed by order size, product mix, or currency effects.

12. Future Work

- Interactive Power BI / dashboard layer for ongoing self-service monitoring of the KPIs in Section 3.
- Sales forecasting using time-series models (e.g. SARIMA or Prophet) informed by the seasonality findings in Section 5.1.
- Customer segmentation via RFM (Recency, Frequency, Monetary) analysis to complement the Pareto findings in Section 5.3.
- Customer Lifetime Value (CLV) modelling to prioritize retention investment by segment.
- Market Basket Analysis to identify cross-sell opportunities, particularly for the underperforming Furniture category.
- Predictive machine learning models (e.g. order-level profit/loss classifier) to flag likely loss-making orders prior to fulfilment.

13. Conclusion

This project demonstrates a complete, reproducible data analytics workflow — from raw transactional data to statistically grounded, prioritized business recommendations. Technology stands out as the company's primary growth driver, while Furniture, high-discount transactions, and a small set of regional markets represent the clearest areas for strategic intervention.

The most important structural finding is that discount governance — not category mix alone — is the dominant controllable driver of profitability in this dataset. Implementing the pricing, product, and regional recommendations in Sections 9–10, and validating each through the proposed controlled experiments, offers a realistic path to improving profitability while preserving the sales growth trajectory observed across 2011–2014.

Prepared by Behnia Shaygan — Data Analytics Portfolio Project